

Tutorial 9

NWEN241

Process Management System Calls

Kirita-Rose Escott
kirita-rose.escott@ecs.vuw.ac.nz

Content

- Process management system calls

- `fork()`

- `exec()`

} Defined in `unistd.h`

- `wait()`

} Defined in `sys/wait.h`

- `exit()`

} Defined in `stdlib.h`

Process creation with **fork()**



- A process calling **fork()** spawns a child process.
- After a successful **fork()** call, two copies of the original code will be running.
 - Parent process – *return value* of `fork()` \rightarrow *child PID*.
 - New child process – *return value* of `fork()` \rightarrow 0.

Example: fork1.c


```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(void) {
    pid_t pid;

    printf("fork test\n");
    if ((pid = fork ()) < 0) {
        printf(" fork error\n");  }
    else if (pid == 0) { /* child */
        printf(" at child process\n");  }
    else { /* parent */
        printf(" at parent process\n");
        wait(NULL);  }
    exit(0);
}
```

Example: fork2.c

```
printf("fork test\n");  
    if ((pid = fork ()) < 0) {  
        printf(" fork error\n");  
    } else if (pid == 0) { /* child */  
        printf(" at child process\n");  
    } else { /* parent */  
        printf(" at parent process\n");  
        wait(NULL);  
    }  
    printf(" where am i?\n");  
    exit(0);
```



How do we know
where we are?

Example: fork3.c

```
printf("fork test\n");
if ((pid = fork ()) < 0) {
    printf(" fork error\n");  }
else if (pid == 0) { /* child */
    printf(" at child process, my pid is: %d\n", (int) getpid());}
else { /* parent */
    printf(" at parent process, my pid is: %d\n", (int) getpid());
    wait(NULL);  }
printf(" i am here with pid : %d\n", (int) getpid());
exit(0);
```

Example: fork4.c

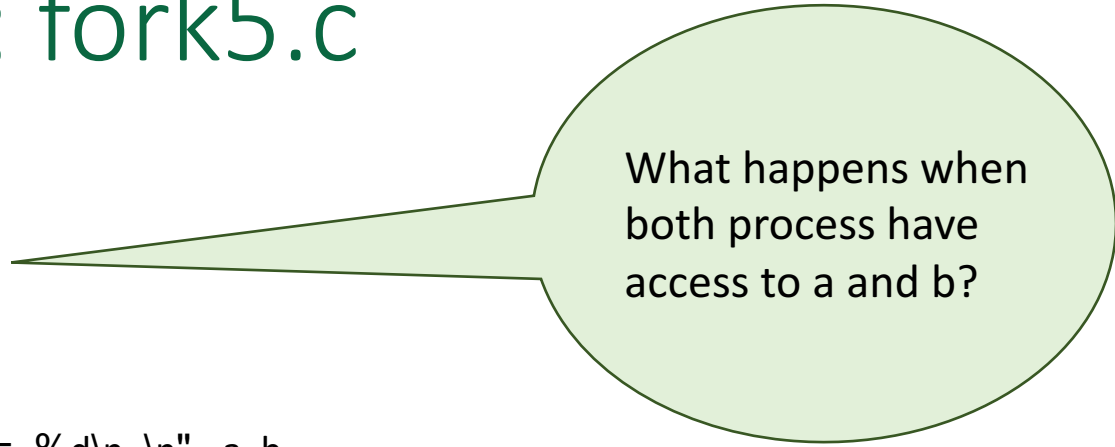
```
#include<stdio.h>
#include<stdlib.h>
#include<unistd.h>

int main(void){
    FILE *fp = fopen("xxx.txt", "a");
    pid_t p = fork();

    if(p < 0) { /* Failed */
        exit(0);
    } else if(p == 0) { /* Child */
        fputs("child\n", fp);
    } else { /* Parent */
        fputs("parent\n", fp);
    }
}
```

Example: fork5.c

```
int main(void) {  
    pid_t pid;  
    int a = 2; int b = 4;  
  
    printf("fork test\n");  
    printf("a = %d and b = %d\n \n", a, b)  
  
    if ((pid = fork ()) < 0) {  
        printf(" fork error\n");  
    } else if (pid == 0) { /* child */  
        a = a * 2;          b = b * 2;  
        printf(" at child process, a = %d and b = %d\n", a, b);  
    } else { /* parent */  
        a = a + 5;          b = b + 5;  
        printf(" at parent process, a = %d and b = %d\n", a, b);  
        wait(NULL);  
    }  
    exit(0);  
}
```



What happens when
both process have
access to a and b?

wait() Call System

- Forces the parent to suspend execution, i.e. wait for its children or a specific child to die (*terminate* is more appropriate terminology, but a bit less common).

pid_t wait(int *status);

- The status, if not NULL, stores exit information of the child, which can be analyzed by the parent.
- The return value is:
 - PID of the exited process, if no error
 - (-1) if an error has happened

exit() System Call

- Gracefully terminates process execution, meaning it does clean up and release of resources, and puts the process into the **zombie** state
→ *terminated **but** still waiting for parent process to read its exit status.*
 - By calling **wait()**, the parent cleans up all its zombie children.
- **exit()** specifies a return value from the program, which a parent process might want to examine as well as status of the dead process.

Example: wait1.c

```
int main(void){
    FILE *fp = fopen("xxx.txt", "a");
    pid_t p = fork();

    if(p < 0) { /* Failed */
        exit(0);
    } else if(p == 0) { /* Child */
        fputs("child says: hi\n", fp);
        fputs("child says: I am\n", fp);
        fputs("child says: the child\n", fp);
    } else { /* Parent */
        wait(NULL);
        fputs("parent says: hello\n", fp);
    }
}
```

What happens when we don't exit the child process? – Example: fork6.c

```
int main(void) {
    pid_t pid;

    printf("fork test\n");
    if ((pid = fork ()) < 0) {
        printf(" fork error\n");
    } else if (pid == 0) { /* child */
        printf(" at child process\n");
        while(1);
    } else { /* parent */
        printf("parent: wait for child\n");
        wait(NULL);
        printf("parent: child complete\n");
    }
    exit(0);
}
```

More about `wait()` and `exit()`

- Should not interpret the status value of system call `wait(&status)` literally. If `&status` is not NULL, `wait()` stores status information in the `int` to which it points.
- Value returned by `exit(&status)` is moved to 2nd byte and 1st (lowest) byte is used to store the status information.

- In previous example:

```
scanf(" %d", &rv); // if value of x is entered
```

```
...
```

```
exit(rv);
```

```
...
```

```
wait(&rv); // the rv contents will be x left shift
           // by 8 bits and additional status
           // written into lowest 8 bit
```



Example: waitexit.c

```
#include <stdio.h>
#include <stdlib.h>

main()
{
    int pid; int rv;

    pid=fork();
    switch(pid){
        case -1:
            printf("Error -- Something went wrong with fork()\n");
            exit(1); // parent exits
        case 0:
            printf("CHILD: This is the child process!\n");
            printf("CHILD: My PID is %d\n", getpid());
            printf("CHILD: My parent's PID is %d\n", getppid());
            printf("CHILD: Enter my exit status: ");
            scanf(" %d", &rv);
            printf("CHILD: I'm outta here!\n");
            exit(rv);
        default:
            printf("PARENT: This is the parent process!\n");
            printf("PARENT: My PID is %d\n", getpid());
            printf("PARENT: My child's PID is %d\n", pid);
            printf("PARENT: I'm now waiting for my child to exit()...\n");
            wait(&rv);
            printf("PARENT: My child's exit status is 0x%.8X\n", rv);
            printf("PARENT: I'm outta here!\n");
    }
}
```

Process creation with `exec()`

- There is **no** system call specifically by the name **`exec()`**
- By **`exec()`** we usually refer to a family of calls:
 - `int execl(char *path, char *arg, ...);`
 - `int execv(char *path, char *argv[]);`
 - `int execlp(char *path, char *arg, ..., char *envp[]);`
 - `int execve(char *path, char *argv[], char *envp[]);`
 - `int execlp(char *file, char *arg, ...);`
 - `int execvp(char *file, char *argv[]);`
- The various options *l*, *v*, *e*, and *p* mean:
 - *l* : an argument list,
 - *v* : an argument vector,
 - *e* : an environment vector, and
 - *p* : a search path.

Example: execl.c

```
#include <stdio.h>
#include <stdlib.h>
int main(void) {
    int ret;

    printf("Calling execl...\n");
    ret = execl("ls", "ls", NULL);
    printf("Failed execl... ret = %d\n", ret);

    return 0;
}
```



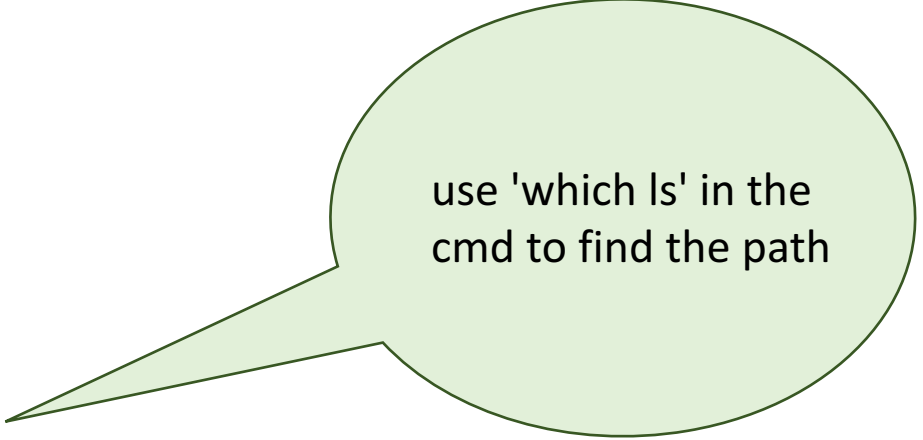
Why does this fail?

Example: corrected execl.c

```
#include <stdio.h>
#include <stdlib.h>
int main(void) {
    int ret;

    printf("Calling execl...\n");
    ret = execl("/bin/l", "l", NULL);
    printf("Failed execl... ret = %d\n", ret);

    return 0;
}
```



use 'which ls' in the
cmd to find the path

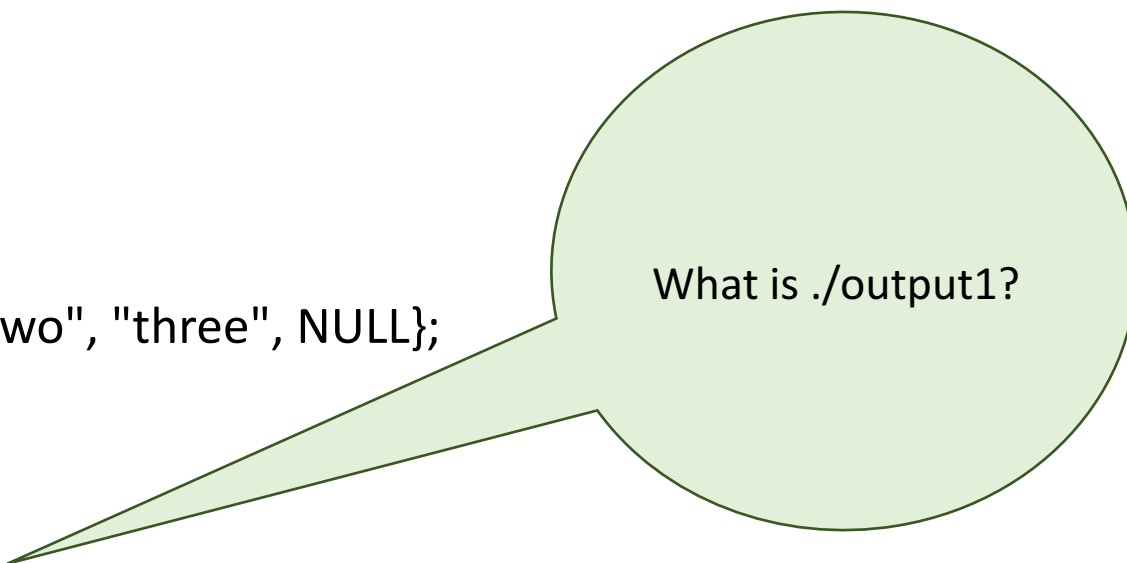
Example: execv.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
int main(void) {
    int ret;
    char *args[] = {"hello", "one", "two", "three", NULL};

    printf("Calling execv...\n");
    ret = execv("./output1", args);
    printf("Failed execv... ret = %d\n", ret);

    exit(0);
}
```



What is ./output1?

Example: output1.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]) {
    int i;
    printf("ARGUMENTS:\n");

    for (i=0; i<argc; i++) {
        printf("%s\n", argv[i]);
    }
    return 0;
}
```

Example: execve.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(void) {
    int ret;
    char *args[] = {"hello", "one", "two", "three", NULL};
    char *envp[] = {"some", "environment", "variables", "ok", NULL};

    printf("Calling execve..\n");
    ret = execv("./output2", args, envp);
    printf("Failed execve... ret = %d\n", ret);

    exit(0);
}
```

Example: output2.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]) {
    int i, j;

    printf("ARGUMENTS:\n");

    for (i=0; i<argc; i++) {
        printf("%s\n", argv[i]);
    }

    printf("ENVIRONMENT VARIABLES:\n");

    for (j=0; j<argc; j++) {
        printf("%s\n", envp[j]);
    }
    return 0;
}
```

Example: execlp.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(void) {
    pid_t pid;

    printf("fork test\n");
    if ((pid = fork ()) < 0) {
        printf(" fork error\n");
    } else if (pid == 0) { /* child */
        printf("child: execute ls command\n");
        execlp("/bin/ls", "ls", NULL);
    } else { /* parent */
        printf("parent: wait for child\n");
        wait(NULL);
        printf("parent: child complete\n");
    }
    exit(0);
}
```