

Week 9 Lecture 1

**NWEN 241**

**Systems Programming**

Alvin C. Valera

`alvin.valera@ecs.vuw.ac.nz`

# Content

## **Low-level Systems Programming**

### Lecture 1:

- Conditional inclusion/compilation
- Standard integer types
- Bit-wise operators

### Lecture 2:

- Bit-fields
- Memory alignment
- Structure padding and packing

# Conditional Inclusion / Compilation

## Recall from Week 1:

- Pre-processor directives (begins with **#**) are instructions to the pre-processor for performing
  - File inclusion
  - Macro substitution
  - *Conditional inclusion*

## Conditional Inclusion (Conditional Compilation):

Capability provided by the preprocessor for selecting lines of code that will be compiled and ignored

# Conditional Inclusion Directives

- Six directives can be used to control conditional compilation

- Beginning of block:

`#if``#ifdef``#ifndef`

- Optional, alternative block:

`#else``#elif`

- End of block:

`#endif`

# Conditional Inclusion Operator

- In addition to the six directives, an operator is also available

`defined name`

`defined (name)`

- Evaluates to 1 if *name* is defined, otherwise it evaluates to 0

# Example using #ifdef and #endif

```
#ifdef MACRO

/**
 * Code here will be compiled
 * if MACRO is defined previously
 */

#endif /* MACRO */
```

# Example using #ifndef and #endif

```
#ifndef MACRO

/**
 * Code here will be compiled
 * if MACRO is not defined previously
 */

#endif /* MACRO */
```

# Example using #if, #elif, #else and #endif

```
#if defined(linux)

/* Code here will be compiled if linux is defined */

#elif defined(_WIN32)

/* Code here will be compiled if _WIN32 is defined */

#else

/* Code here will be compiled if neither linux or _WIN32
   is defined */

#endif
```



# Where to define a macro

- Within a header or source file
  - Using `#define` directive
- Within a `Makefile`
  - To be discussed in “Writing Large Program” Tutorial on Friday
- As a command-line option to `gcc` or `g++`
  - `gcc -DMACRO_NAME source.c`: `MACRO_NAME` will be defined in `source.c`
  - `g++ -DMACRO_NAME source.cc`: `MACRO_NAME` will be defined in `source.cc`

# Example

hello.c

```
#include <stdio.h>

int main(void)
{
#ifdef HELLO
    printf("hello\n");
#else
    printf("world\n");
#endif
    return 0;
}
```

```
$ gcc -DHELLO hello.c
$ ./a.out
hello

$ gcc hello.c
$ ./a.out
world
```

# Recall: Basic C/C++ Data Types

Data Type	Size (bytes)	
<del>boolean</del>	<del>1</del>	Integral types
<del>byte</del>	<del>1</del>	
char	<del>2</del> 1	
short (short int)	<del>2</del> Machine-dependent	
int	<del>4</del> Machine-dependent	
long (long int)	<del>8</del> Machine-dependent	
<del>long long (long long int)</del>	<del>Machine-dependent</del>	Float types
float	<del>4</del> Machine-dependent	
double	<del>8</del> Machine-dependent	
<del>long double</del>	<del>10</del>	

# Integer Types

- A major problem in systems programming is the uncertainty of integer types
  - Size of an integer depends on the CPU architecture
- There are situations where you will need to specify the size precisely

Fortunately, C99 defines integer types with precise sizes

# Fixed Size Integer Types


- Defined in `stdint.h` (C) and `cstdint` (C++)

Type	Sign and size
<code>int8_t</code>	8-bit signed integer
<code>int16_t</code>	16-bit signed integer
<code>int32_t</code>	32-bit signed integer
<code>int64_t</code>	64-bit signed integer
<code>uint8_t</code>	8-bit unsigned integer
<code>uint16_t</code>	16-bit unsigned integer
<code>uint32_t</code>	32-bit unsigned integer
<code>uint64_t</code>	64-bit unsigned integer


# Binary Number Representation

- How are numbers actually stored in variables?
  - Numbers are stored using binary representation
  - See <https://www.bottomupcs.com/chapter01.xhtml> for details

```
char c = 'A'; //65
```



```
uint8_t b = 255;
```



# Binary Number Representation

How are negative numbers represented?

- Use **two's complement**

## **Two's Complement Operation:**

1. Obtain binary representation disregarding the sign
2. If sign is positive, done
3. Otherwise:
  1. Invert all bits of the binary representation
  2. Add 1 to the binary representation

# Binary Number Representation

- Example:

```
int8_t a = 37;
```



0	0	1	0	0	1	0	1
$2^7$	$2^6$	$2^5$	$2^4$	$2^3$	$2^2$	$2^1$	$2^0$

```
int8_t b = -37;
```



Invert all bits:

0	0	1	0	0	1	0	1
1	1	0	1	1	0	1	0


Add 1:

1	1	0	1	1	0	1	1
---	---	---	---	---	---	---	---



# Binary Number Representation

- When you get the two's complement of a negative number, you will obtain the positive number



# When to use unsigned types?

- Passing and comparing signed and unsigned values is a common pitfall in C/C++ programming

```
#include <stdio.h>
int array[] = {1, 2, 3};
void dump(unsigned int len)
{
    for(unsigned int i=0; i<len; i++)
        printf("array[%d]: %d\n", i, array[i]);
}
int main(void)
{
    dump(-1);
    return 0;
}
```

Will be treated as a very large positive integer by dump()

# When to use unsigned types?

- If possible, avoid using unsigned types
- Areas where unsigned types are used:
  - When dealing with bit values
  - Performing bit-level operations

# Endianness

- Consider a 32-bit (4-byte) integer:

```
uint32_t a = 305419896;
```

```
0001001000110100010101100111000
```

- How is it actually stored in computer memory?

There are two ways of storing multi-byte numbers in memory:

- **Big Endian**
- **Little Endian**

# Endianness

00010010 00110100 01010110 01111000

Most significant  
byte (MSB)

Least significant  
byte (LSB)

- Big Endian – Store MSB in lower address
- Little Endian – Store LSB in lower address

# Endianness

- Consider a 32-bit (4-byte) integer:

```
uint32_t a = 305419896;
```

00010010 00110100 01010110 01111000

- Suppose a is at address 100

Big Endian:

100 00010010

101 00110100

102 01010110

103 01111000

Little Endian:

100 01111000

101 01010110

102 00110100

103 00010010

# Bit-wise Operators

## & – AND

- Result is **1** if both operand bits are **1**

## | – OR

- Result is **1** if either operand bit is **1**

## ^ – Exclusive OR

- Result is **1** if operand bits are different

## ~ – Complement

- Each bit is reversed

## << – Shift left

- Multiply by 2

## >> – Shift right

- Divide by 2

# Example: AND

```
uint8_t a = 37;  
uint8_t b = 98;  
uint8_t c = a & b;
```

0	0	1	0	0	1	0	1
---	---	---	---	---	---	---	---

0	1	1	0	0	0	1	0
---	---	---	---	---	---	---	---

0	0	1	0	0	0	0	0
---	---	---	---	---	---	---	---



# Example: OR

```
uint8_t a = 37;
```

```
uint8_t b = 98;
```

```
uint8_t c = a | b;
```

0	0	1	0	0	1	0	1
---	---	---	---	---	---	---	---

0	1	1	0	0	0	1	0
---	---	---	---	---	---	---	---

0	1	1	0	0	1	1	1
---	---	---	---	---	---	---	---

# Example: Exclusive OR

```
uint8_t a = 37;  
uint8_t b = 98;  
uint8_t c = a | b;
```

0	0	1	0	0	1	0	1
---	---	---	---	---	---	---	---

0	1	1	0	0	0	1	0
---	---	---	---	---	---	---	---

0	1	0	0	0	1	1	1
---	---	---	---	---	---	---	---

# Example: Complement

```
uint8_t a = 37;
```

```
uint8_t b = ~a;
```

0	0	1	0	0	1	0	1
---	---	---	---	---	---	---	---

1	1	0	1	1	0	1	0
---	---	---	---	---	---	---	---

# Example: Shift Left

```
uint8_t a = 37;  
uint8_t b = a << 1;  
uint8_t c = a << 2;
```

0	0	1	0	0	1	0	1
---	---	---	---	---	---	---	---

0	1	0	0	1	0	1	0
---	---	---	---	---	---	---	---

1	0	0	1	0	1	0	0
---	---	---	---	---	---	---	---

What is b and c in decimal?

# Example: Shift Right

```
uint8_t a = 37;  
uint8_t b = a >> 1;  
uint8_t c = a >> 2;
```

0	0	1	0	0	1	0	1
---	---	---	---	---	---	---	---

0	0	0	1	0	0	1	0
---	---	---	---	---	---	---	---

0	0	0	0	1	0	0	1
---	---	---	---	---	---	---	---

What is b and c in decimal?

# Content

## **Next Lecture**

- Bit-fields
- Memory alignment
- Structure padding and packing