

Week 9 Lecture 2

**NWEN 241**

**Systems Programming**

Alvin C. Valera

`alvin.valera@ecs.vuw.ac.nz`

# Content

## **Low-level Systems Programming**

### Lecture 1:

- Conditional inclusion/compilation
- Standard integer types
- Bit-wise operators

### Lecture 2:

- Bit-fields
- Memory alignment
- Structure padding and packing

# Accessing bits or groups of bits

## Two Approaches:

### Traditional C:

Use #define macro in tandem with bitwise operators

### Modern:

Use bit-fields

# Traditional approach: selecting bits

Suppose

```
uint8_t a = 37;
```


7	6	5	4	3	2	1	0
0	0	1	0	0	1	0	1

How to select / read the lowest 3 bits (bits 0, 1 and 2)?

	0	0	1	0	0	1	0	1
&	0	0	0	0	0	1	1	1
	0	0	0	0	0	1	0	1

Bit mask: indicates bits to be selected

# Traditional approach: selecting bits



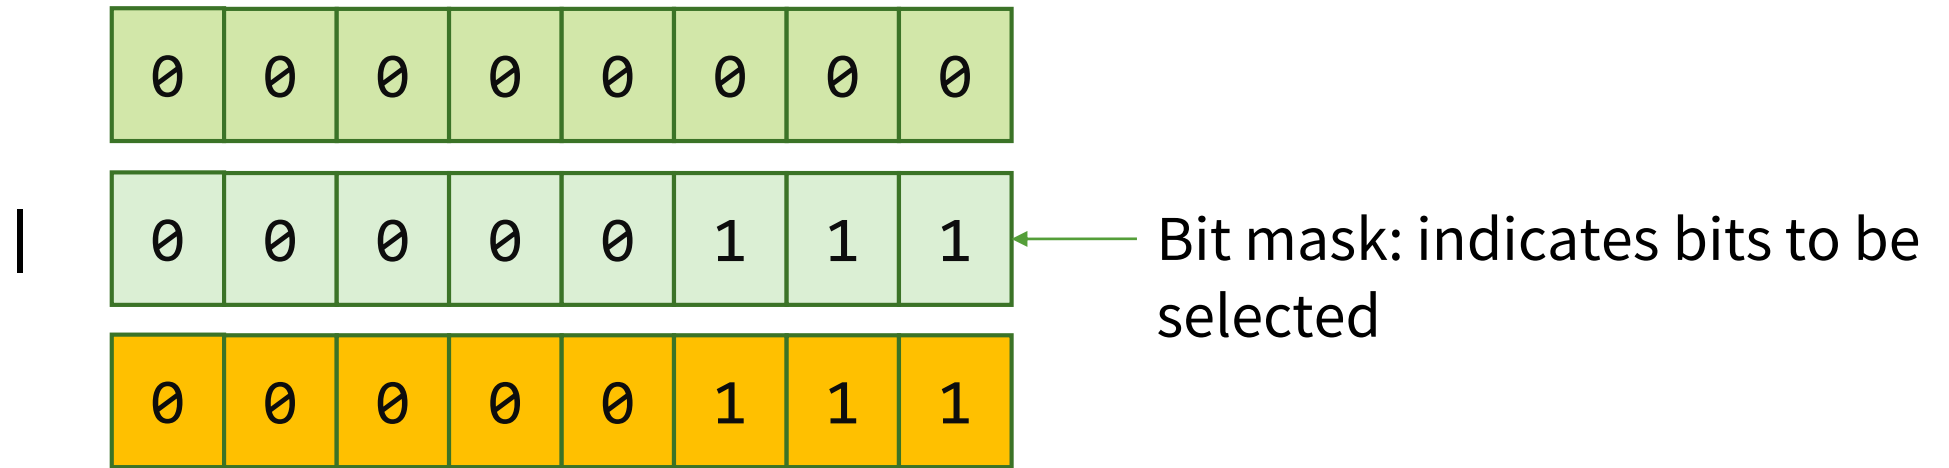
```
uint8_t a = 37;  
uint8_t mask = 0x07; // binary 00000111  
uint8_t b = a & mask;
```

# Binary to hex

Binary	Hex
0000	0
0001	1
0010	2
0011	3
0100	4
0101	5
0110	6
0111	7


Binary	Hex
1000	8
1001	9
1010	A
1011	B
1100	C
1101	D
1110	E
1111	F

# Traditional approach: setting bits



```
uint8_t a = 0;  
uint8_t mask = 0x07; // binary 00000111  
uint8_t b = a | mask;
```

# Traditional approach: clearing bits



```
uint8_t a = 255;  
uint8_t mask = 0x07; // binary 00000111  
uint8_t b = a & ~mask;
```



# Traditional approach example

- Serial port line register is an 8-bit read-only register:

7	6	5	4	3	2	1	0
---	---	---	---	---	---	---	---

Bit 7: FIFO Error

Bit 6: Empty Data Holding Registers

Bit 5: Empty Transmitter Holding Registers

Bit 4: Break Interrupt

Bit 3: Framing Error

Bit 2: Parity Error

Bit 1: Overrun Error

Bit 0: Data Ready

- When bit is set (1), the condition exists

# Traditional approach example

```
#define FIFO_ERROR      0x80 // bit 7 mask
#define EMPTY_DHR      0x40 // bit 6 mask
#define EMPTY_THR      0x20 // bit 5 mask
#define BREAK_INTR     0x10 // bit 4 mask
#define FRAMING_ERROR  0x08 // bit 3 mask
#define PARITY_ERROR    0x04 // bit 2 mask
#define OVERRUN_ERROR   0x02 // bit 1 mask
#define DATA_READY     0x01 // bit 0 mask
```

```
void foo(void)
{
    uint8_t reg = read_line_register();
    if (reg & DATA_READY) {
        // Data is ready
    }
}
```

# Traditional approach

- Traditional approach works very well for bit-wise access
- For multi-bit access, use **bit-fields**

A **bit-field** is a data structure that allows access and/or operation of individual bits or group of bits of a word

# Bit-fields in C/C++

- Declared as a struct
  - Each member is a bit-field within a word
  - Accessed like members of a struct
  - Fields may be named or un-named

```
struct structure_tag {  
    typeA memberA1 : bit_widthA1;  
    typeA memberA2 : bit_widthA2;  
    ...  
    typeB memberB1 : bit_widthB1;  
    typeB memberB2 : bit_widthB2;  
    ...  
} variable_list;
```

# Bit-fields example (serial port line reg.)

```
struct sp_line_reg {  
    uint8_t fifo_error    : 1;  
    uint8_t empty_dhr     : 1;  
    uint8_t empty_thr     : 1;  
    uint8_t break_intr    : 1;  
    uint8_t framing_error : 1;  
    uint8_t parity_error  : 1;  
    uint8_t overrun_error : 1;  
    uint8_t data_ready    : 1;  
};  
  
void foo(void)  
{  
    struct sp_line_reg reg = read_line_register();  
    if (reg.data_ready) {  
        // Data is ready  
    }  
}
```

# Bit-fields example (serial port line reg.)

```
struct sp_line_reg {  
    uint8_t fifo_error    : 1;  
    uint8_t empty_dhr     : 1;  
    uint8_t empty_thr     : 1;  
    uint8_t break_intr    : 1;  
    uint8_t framing_error : 1;  
    uint8_t parity_error  : 1;  
    uint8_t overrun_error : 1;  
    uint8_t data_ready    : 1;  
};
```



```
struct sp_line_reg {  
    uint8_t fifo_error    : 1,  
           empty_dhr      : 1,  
           empty_thr      : 1,  
           break_intr     : 1,  
           framing_error  : 1,  
           parity_error   : 1,  
           overrun_error  : 1,  
           data_ready     : 1;  
};
```

# Bit-fields example (vs traditional)

```
struct sp_line_reg {
    uint8_t fifo_error    : 1;
    uint8_t empty_dhr     : 1;
    uint8_t empty_thr     : 1;
    uint8_t break_intr    : 1;
    uint8_t framing_error : 1;
    uint8_t parity_error  : 1;
    uint8_t overrun_error : 1;
    uint8_t data_ready    : 1;
};

void foo(void)
{
    struct sp_line_reg reg =
        read_line_register();
    if (reg.data_ready) {
        // Data is ready
    }
}
```

```
#define FIFO_ERROR      0x80 // bit 7 mask
#define EMPTY_DHR      0x40 // bit 6 mask
#define EMPTY_THR      0x20 // bit 5 mask
#define BREAK_INTR     0x10 // bit 4 mask
#define FRAMING_ERROR   0x08 // bit 3 mask
#define PARITY_ERROR    0x04 // bit 2 mask
#define OVERRUN_ERROR   0x02 // bit 1 mask
#define DATA_READY     0x01 // bit 0 mask

void foo(void)
{
    uint8_t reg =
        read_line_register();
    if (reg & DATA_READY) {
        // Data is ready
    }
}
```

# Problems with bit-fields

- The actual arrangement of the bits on memory depends on the compiler and/or “endian-ness” of the CPU
- **Bit-fields are not portable**

```
struct sp_line_reg {  
    uint8_t fifo_error    : 1;  
    uint8_t empty_dhr     : 1;  
    uint8_t empty_thr     : 1;  
    uint8_t break_intr    : 1;  
    uint8_t framing_error : 1;  
    uint8_t parity_error  : 1;  
    uint8_t overrun_error : 1;  
    uint8_t data_ready    : 1;  
};
```

Can be mapped in memory  
this way:





# Problems with bit-fields

- The actual arrangement of the bits on memory depends on the compiler and/or “endian-ness” of the CPU
- **Bit-fields are not portable**

Or possibly this way:

```
struct sp_line_reg {
```

```
    uint8_t fifo_error    : 1;
```

```
    uint8_t empty_dhr     : 1;
```

```
    uint8_t empty_thr     : 1;
```

```
    uint8_t break_intr    : 1;
```

```
    uint8_t framing_error : 1;
```

```
    uint8_t parity_error  : 1;
```

```
    uint8_t overrun_error : 1;
```

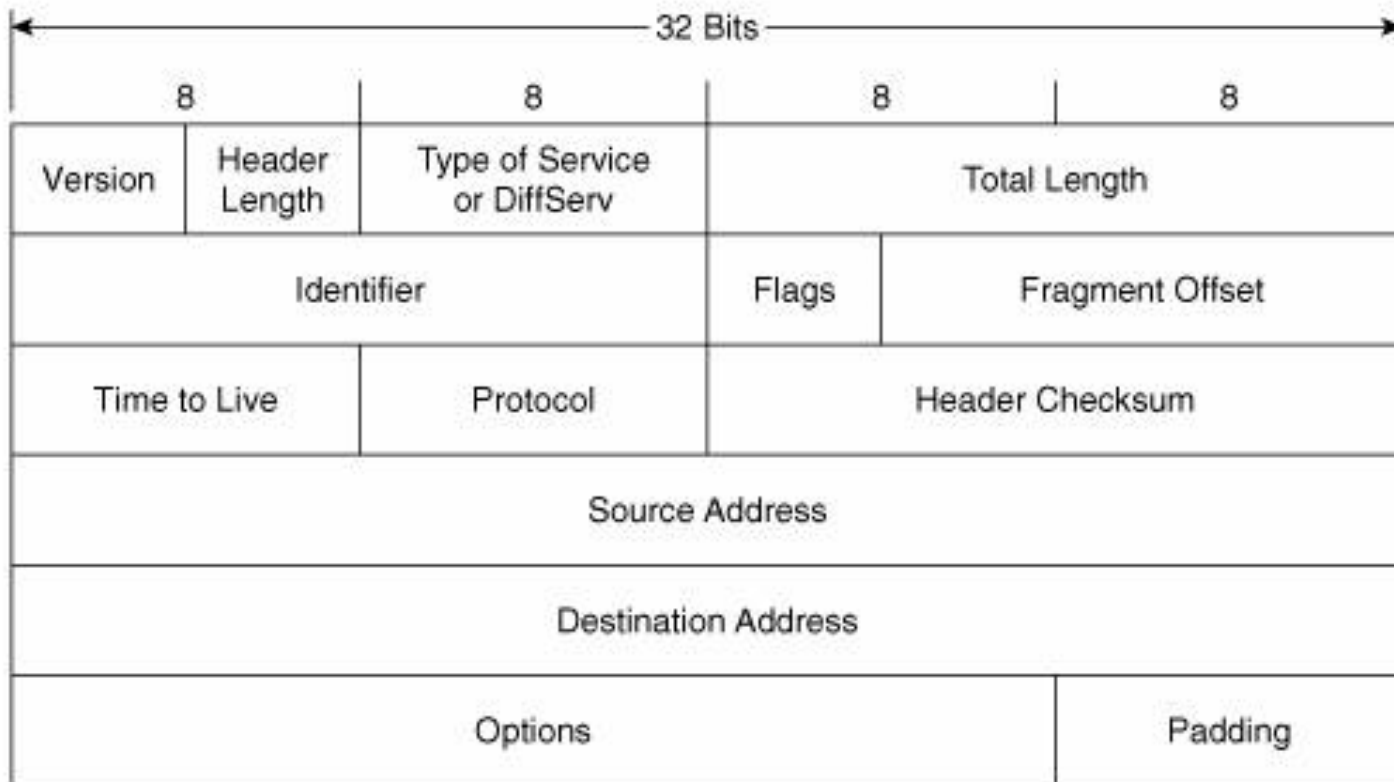
```
    uint8_t data_ready    : 1;
```

```
};
```

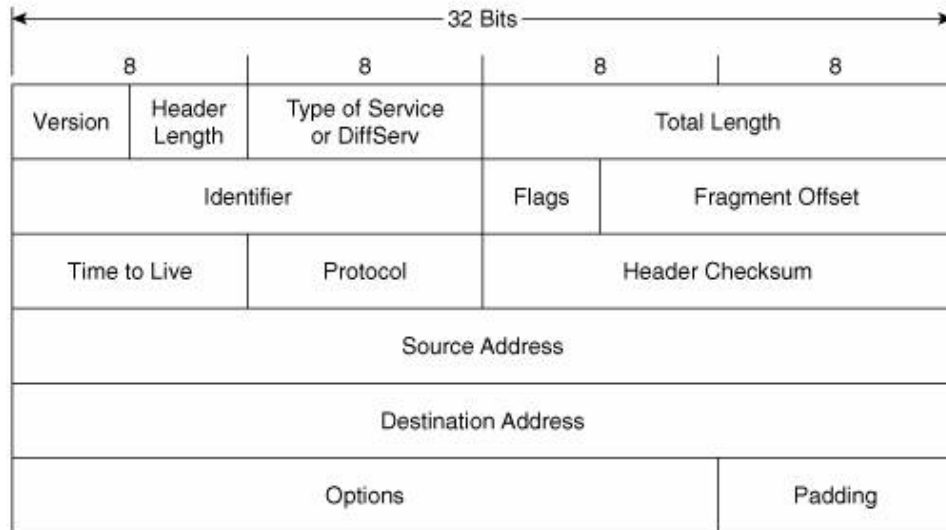


# Bit-fields example: IPv4 packet header

Internet Protocol version 4 packet header format:



# Bit-fields example: IPv4 packet header



```
struct iphdr {
    #if defined(__LITTLE_ENDIAN_BITFIELD)
        uint8_t ihl      :4,
                version:4;
    #elif defined (__BIG_ENDIAN_BITFIELD)
        uint8_t version:4,
                ihl      :4;
    #else
        #error "Please fix <asm/byteorder.h>"
    #endif
    uint8_t  tos;
    uint16_t tot_len;
    uint16_t id;
    uint16_t frag_off;
    uint8_t  ttl;
    uint8_t  protocol;
    uint16_t check;
    uint32_t saddr;
    uint32_t daddr;
};
```

# Bit-fields example: IPv4 packet header

```
void ipv4_receive(void *pkt, uint32_t my_addr)
{
    struct iphdr *iph = (struct iphdr *) pkt;

    if(iph->version != 4) {
        // Incorrect version, return
        return;
    }

    if(iph->daddr == my_addr) {
        // This packet is for me
        // Do stuff to receive it!
    }
}
```

# Memory alignment

- Many computer systems place restrictions on the allowable addresses for the basic data types
  - The address for some type must be a multiple of some value  $k$  (typically 2, 4, or 8)
- Such **alignment restriction** is meant for improving memory access performance

# Linux memory alignment policy on 32-bit x86

- 1-byte data types (e.g., `int8_t`, `uint8_t`, `char`) can have any address
- 2-byte data types (e.g., `int16_t`, `uint16_t`, `short`) must have an address that is a multiple of 2
- Larger data types (e.g., `int32_t`, `int64_t`, `float`, `double`) must have an address that is a multiple of 4

# Memory alignment on structures


- Compiler may need to insert gaps in the field allocation to ensure that each structure element satisfies its alignment requirement
- The entire structure then has to follow a required alignment for its starting address

# Memory alignment on structures

- Consider the following:

```
struct s1 {  
    int32_t a;  
    int8_t b;  
    int32_t c;  
};
```

- Minimally, s1 can be allocated in memory:




- Member c is not aligned!
  - Its address should be a multiple of 4




# Structure padding

- Minimally, s1 can be allocated in memory:




- Member c is not aligned!
  - Its address should be a multiple of 4

- Solution: Compiler will add **padding** to align members



# What else?

- Solution: Compiler will add **padding** to align members



- In addition to padding some of members, the structure itself should begin at an address that is a multiple of 4

# Structure packing

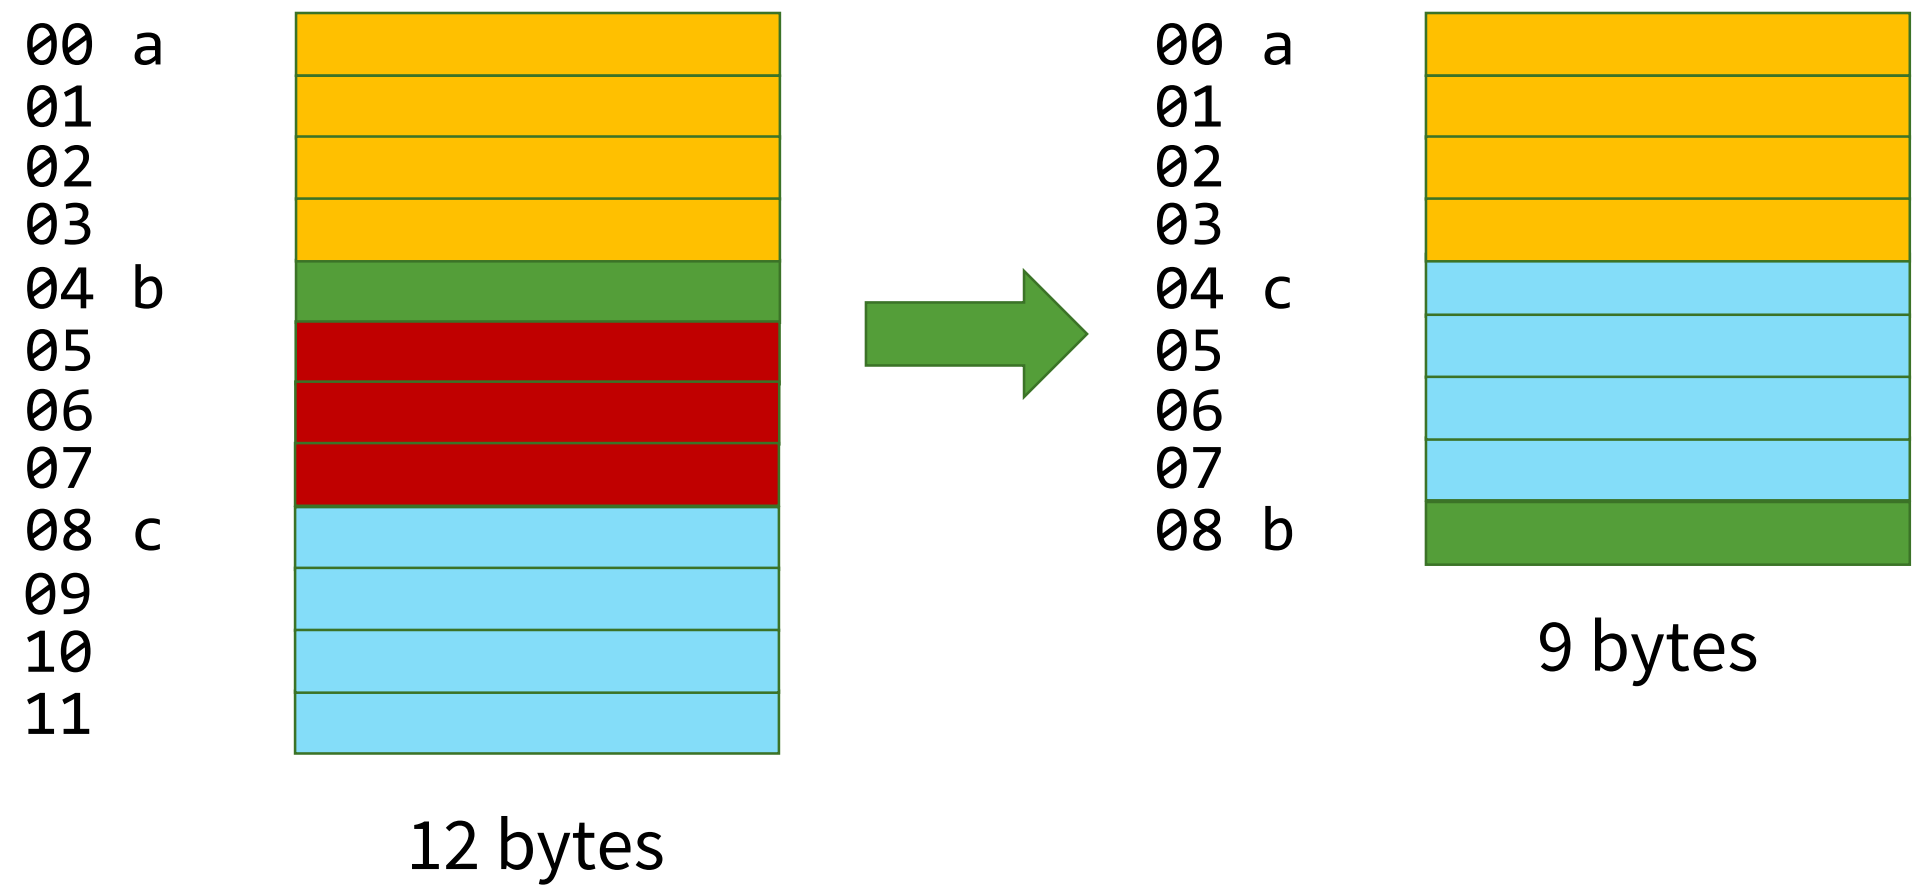
- Packing: re-arranging structure members to reduce waste due to padding for data alignment

```
struct s1 {  
    int32_t a;  
    int8_t b;  
    int32_t c;  
};
```



```
struct s1 {  
    int32_t a;  
    int32_t c;  
    int8_t b;  
};
```


# Structure packing



# Overriding alignment

- GCC and G++ provides an option to override alignment restrictions on structures
  - May not work on all CPUs
  - Will result in slower code

```
struct s1 {  
    int32_t a;  
    int8_t b;  
    int32_t c;  
};
```



```
struct s1 {  
    int32_t a;  
    int32_t c;  
    int8_t b;  
} __attribute__((packed));
```

# Next lecture

- Process management