

Week 11 Lecture 2

NWEN 241

Systems Programming

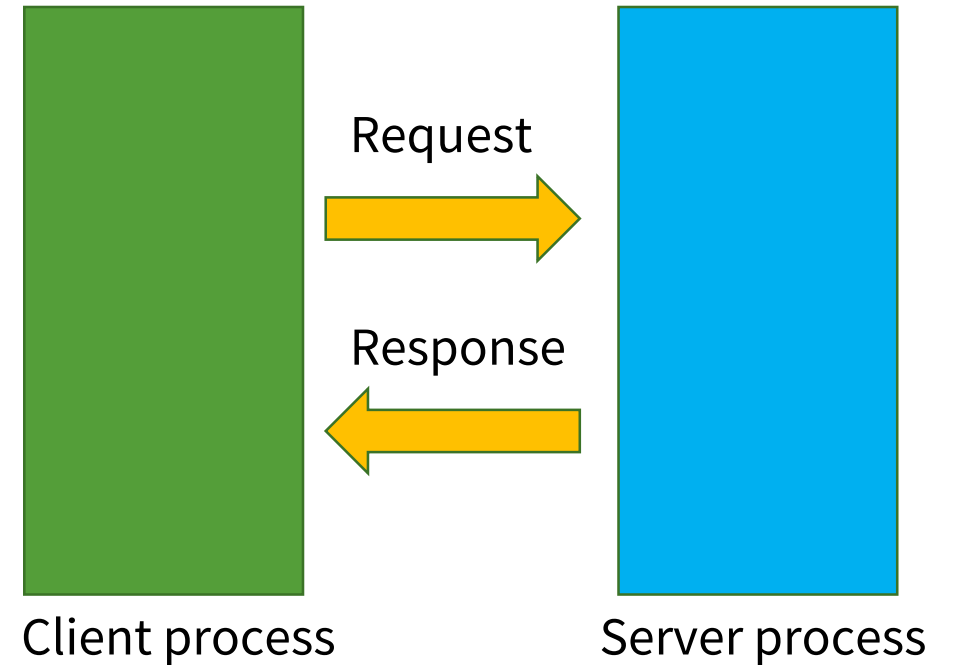
Alvin C. Valera
`alvin.valera@ecs.vuw.ac.nz`

Content

- Socket programming

Recap: Client-server model

- Based on the producer-consumer model of process cooperation
- Client makes the request for some resource or service to the server process
- Server process handles the request and sends the response (result) back to the client



Recap: Client-server model

- Client process needs to know the existence and the address of the server
- However, the server does not need to know the existence or address of the client prior to the connection
- Once a connection is established, both sides can send and receive information

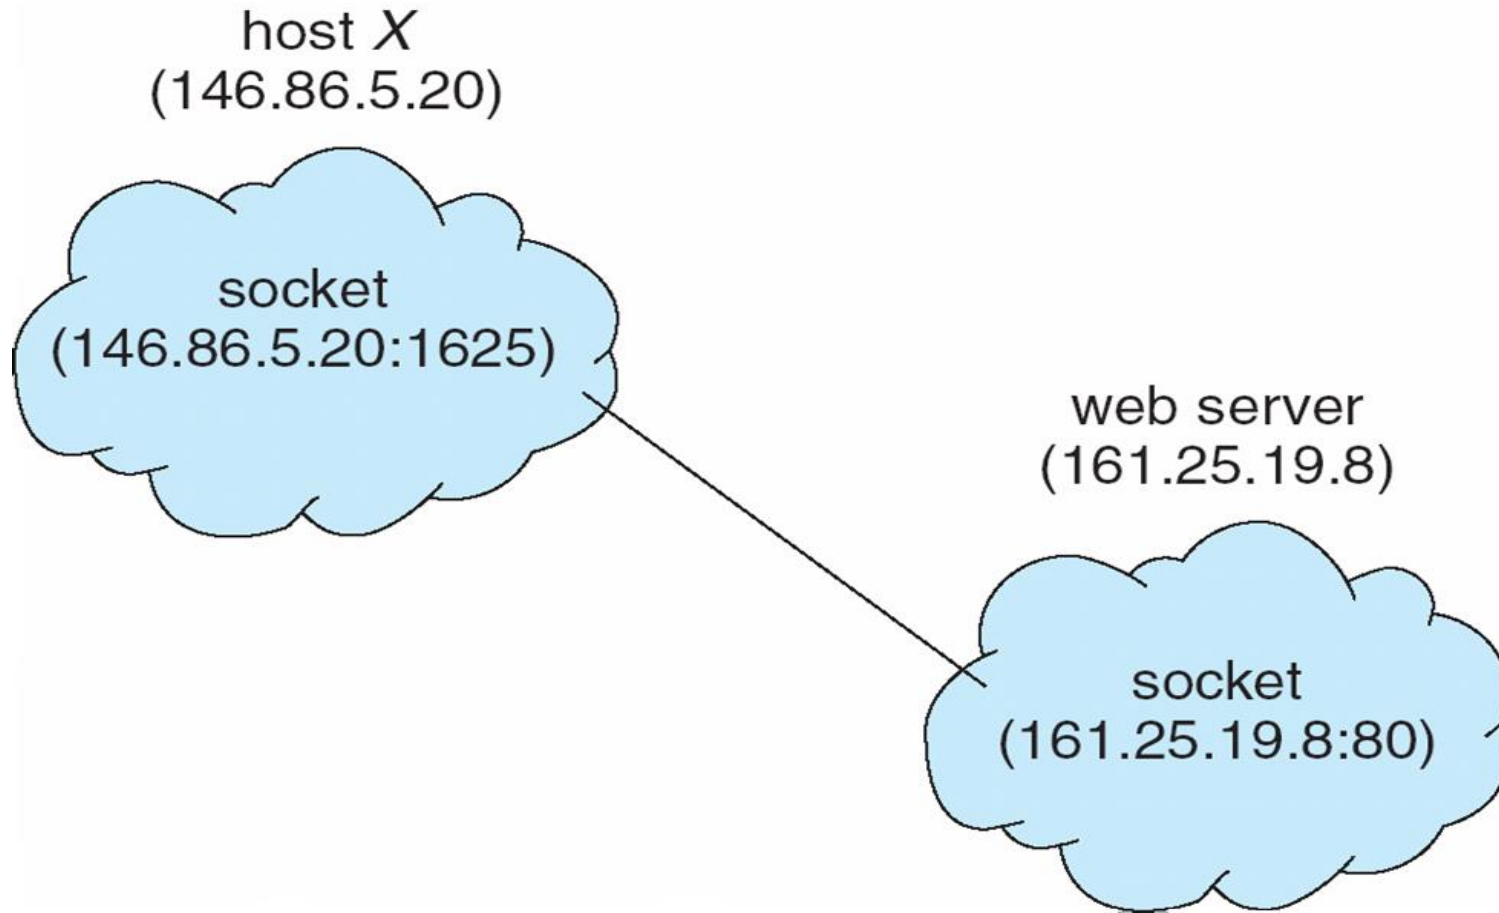
Recap: Client-server communication

- Remote Procedure Calls
- Pipes
- Remote Method Invocation (Java)
- **Sockets**

What is socket?

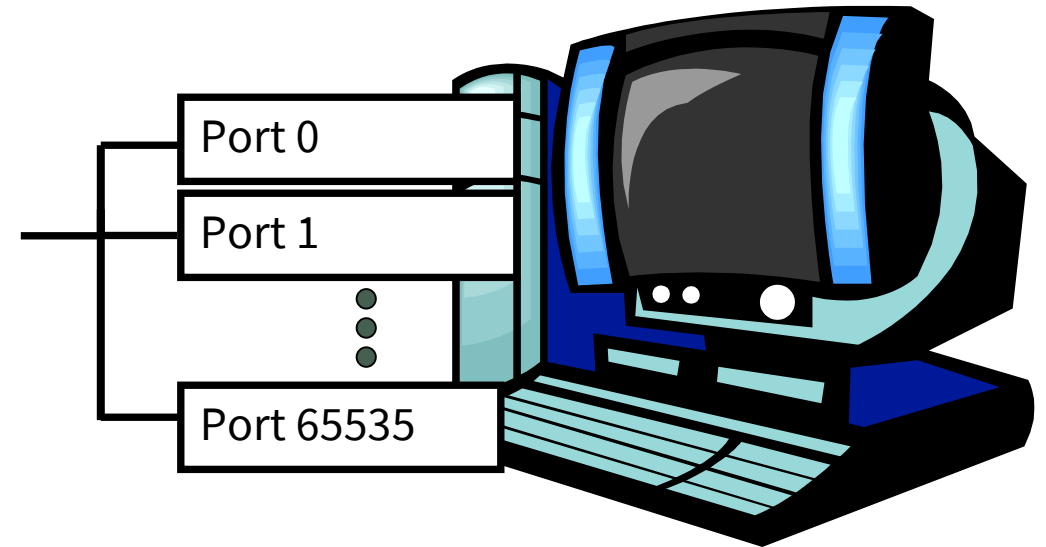
- A **socket** is defined as an endpoint for communication
- Concatenation of **IP address** and **port** – a number included at start of message packet to differentiate network services on a host
- Example:
 - The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**

Socket communication



Port numbers

- Each host has 65,536 ports
- Use of ports 1-1024 requires privileges
- Some ports are reserved for specific apps
 - 20, 21: FTP
 - 23: Telnet
 - 80: HTTP
 - see RFC 1700 (about 2000 ports are reserved)



Sockets as programming interface

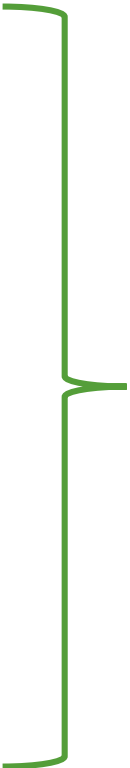
- An interface between application and network
 - The application creates a socket
 - The socket type dictates the style of communication
 - reliable vs. best effort
 - connection-oriented vs. connectionless

Socket types

- SOCK_STREAM
 - a.k.a. **TCP**
 - reliable delivery
 - in-order guaranteed
 - connection-oriented
 - bidirectional
- SOCK_DGRAM
 - a.k.a. **UDP**
 - unreliable delivery
 - no order guarantees
 - no notion of “connection” – app indicates dest. for each packet
 - can send or receive

System calls

- `socket()`
- `bind()`
- `listen()`
- `accept()`
- `connect()`
- `send()` / `sendto()`
- `recv()` / `recvfrom()`



Defined in
`sys/types.h`
`sys/socket.h`

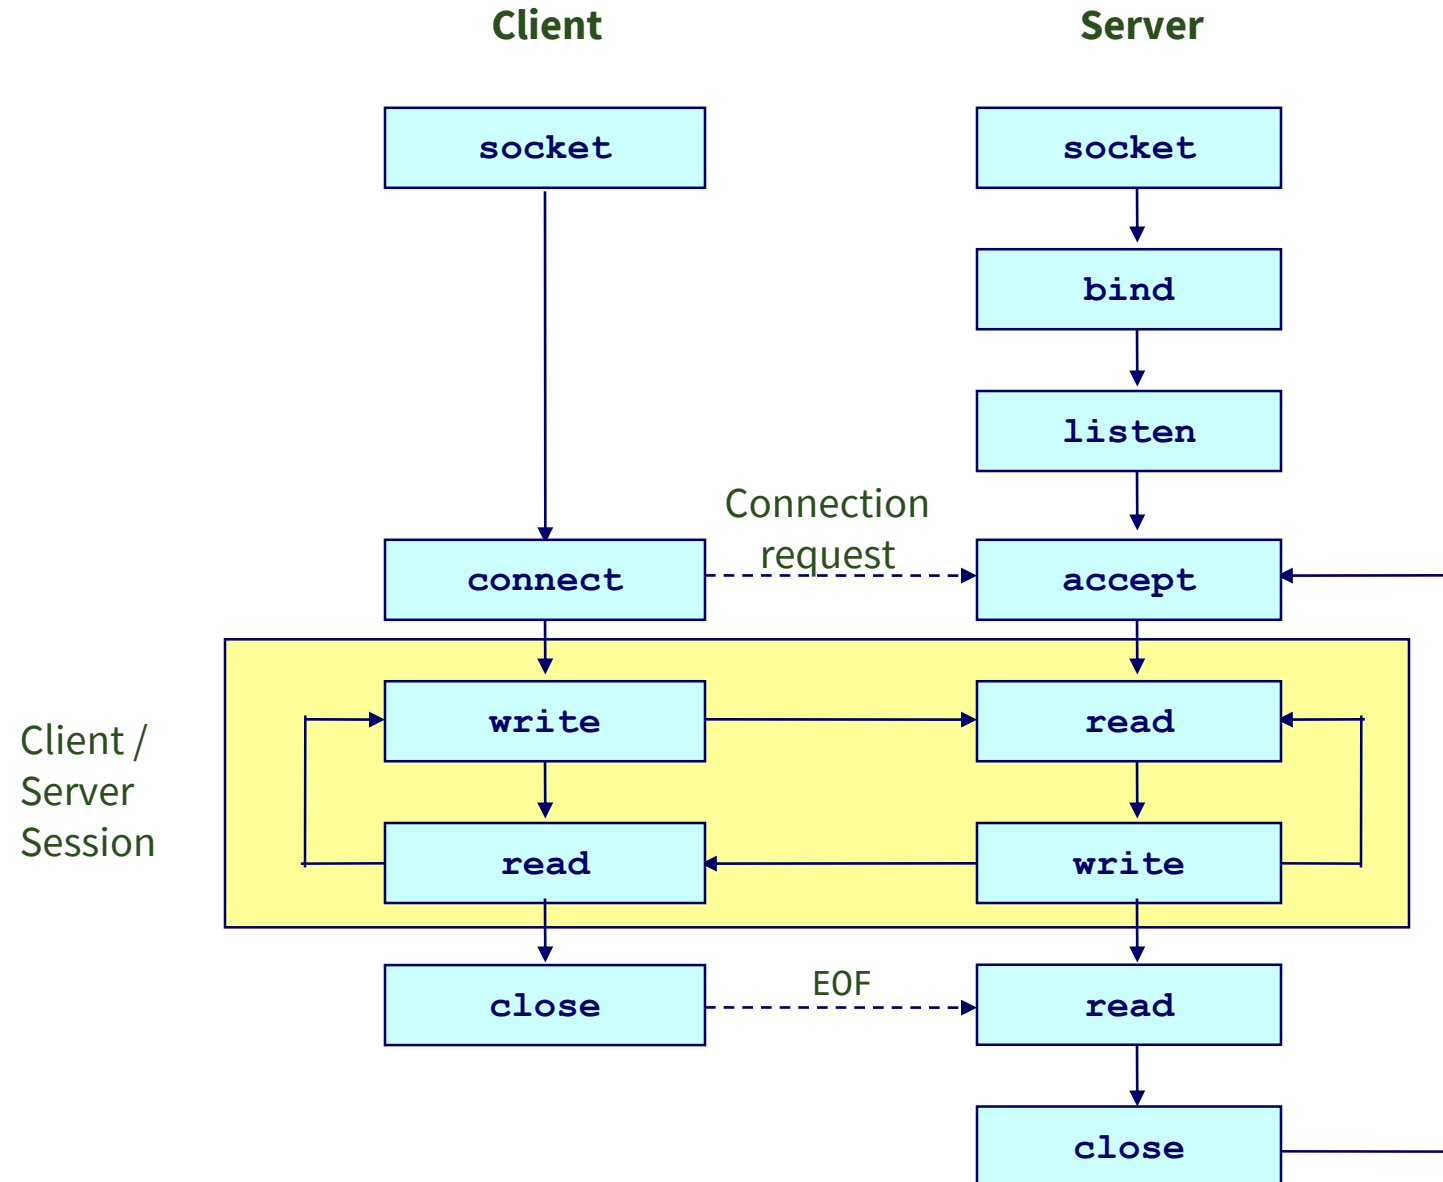
Server overview

- 1) Create a socket with the `socket()` system call
- 2) Bind the socket to an address using the `bind()` system call
- 3) Listen for connections with the `listen()` system call
- 4) Accept a connection with the `accept()` system call
- 5) Send and receive data

Client overview

- 1) Create a socket with the `socket()` system call
- 2) Connect the socket to the address of the server using the `connect()` system call
- 3) Send and receive data

Client-server communication overview



Server: step 1

- Create a socket with the `socket()` system call

```
int socket(int domain, int type, int protocol);
```

- *domain* can either be `AF_INET` (IPv4) or `AF_INET6` (IPv6)
 - *type* can either be `SOCK_STREAM` (TCP) or `SOCK_DGRAM` (UDP)
 - *protocol* specifies the protocol, usually 0.
-
- If successful, returns **socket file descriptor**, otherwise, returns -1

Server: step 1 example

- Create TCP socket

```
int fd = socket(AF_INET, SOCK_STREAM, 0);  
if (fd == -1) {  
    printf("Error creating socket");  
    exit(0);  
}
```

- Create UDP socket

```
int fd = socket(AF_INET, SOCK_DGRAM, 0);  
if (fd == -1) {  
    printf("Error creating socket");  
    exit(0);  
}
```


Server: step 2

- Bind the socket to an address using the `bind()` system call
 - Binding means associating and reserving a port number for use by the socket

```
int bind(int sockfd, const struct sockaddr *addr,  
         socklen_t addrlen);
```

- *sockfd* is the socket file descriptor (returned by `socket()`)
- *addr* is a pointer to the structure `struct sockaddr` which contains the host IP address and port number to bind to
- *addrlen* is the length of what *addr* points to
- If successful, returns 0, otherwise, returns -1

struct sockaddr

- Uses struct sockaddr_in in IPv4

```
struct sockaddr_in {  
    short sin_family;           // e.g. AF_INET  
    unsigned short sin_port;    // port number  
    struct in_addr sin_addr;    // address  
    char sin_zero[8];          // padding to be same size  
                                // as struct sockaddr  
};  
  
struct in_addr {  
    unsigned long s_addr;       // IPv4 address  
};
```

Server: step 2 example

```
int fd = socket(AF_INET, SOCK_STREAM, 0);  
...  
  
struct sockaddr_in addr;  
addr.sin_family = AF_INET;  
addr.sin_addr.s_addr = INADDR_ANY; // any address  
addr.sin_port = htons(12345);  
  
if (bind(fd, (struct sockaddr *)&addr, sizeof(addr)) < 0) {  
    printf("Error binding socket");  
    exit(0);  
}
```

Host and network byte order

- Remember the little-endian and big-endian issue?
- Byte ordering also matters in network communication
 - Host and network may differ in byte ordering
- Functions for converting between host and network byte order:

```
uint32_t htonl(uint32_t hostlong);  
uint16_t htons(uint16_t hostshort);  
uint32_t ntohl(uint32_t netlong);  
uint16_t ntohs(uint16_t netshort);
```

Server: step 3

- Listen for connections with the `listen()` system call

```
int listen(int sockfd, int backLog);
```

- *sockfd* is the socket file descriptor (returned by `socket()`)
- *backLog* is the maximum number of pending connections to allow for this socket
 - `SOMAXCONN` is defined as the number of maximum pending connections allowed by the operating system
- If successful, returns 0, otherwise, returns -1

Server: step 3 example

```
int fd = socket(AF_INET, SOCK_STREAM, 0);  
...  
  
if(listen(fd, SOMAXCONN) < 0) {  
    printf("Error listening for connections");  
    exit(0);  
}
```

Server: step 4

- Accept a connection with the `accept()` system call

```
int accept(int sockfd, struct sockaddr *addr,  
           socklen_t *addrlen);
```

- `sockfd` is the socket file descriptor (returned by `socket()`)
- `addr` is a pointer to the structure `struct sockaddr` which will contain the details of the peer socket
- `addrlen` is a pointer to the length of what `addr` points to
- If successful, returns non-negative **socket file descriptor**, otherwise, returns -1

Server: step 4 example

```
int fd = socket(AF_INET, SOCK_STREAM, 0);  
...  
struct sockaddr_in client_addr;  
int addrlen = sizeof(client_addr);  
  
int client_fd = accept(fd, (struct sockaddr *)&client_addr,  
                      (socklen_t*)&addrlen);  
  
if(client_fd < 0) {  
    printf("Error accepting connection");  
    exit(0);  
}
```


Server: step 5

- **Send** and receive data

```
ssize_t send(int sockfd, const void *buf, size_t len, int flags);
```

```
ssize_t sendto(int sockfd, const void *buf, size_t len, int flags,  
               const struct sockaddr *dest_addr, socklen_t addrlen);
```

- *sockfd* is the socket file descriptor (returned by socket())
- *buf* is a pointer to buffer to be sent
- *len* is the length of buffer to be sent
- *flags* is bitwise OR of zero or more options
- *dest_addr* is a pointer to the structure struct `sockaddr` which will contain the details of the peer socket
- *addrlen* is a pointer to the length of what *dest_addr* points to
- If successful, returns number of characters sent, otherwise, returns -1

Server: step 5

- **Send** and receive data

```
ssize_t send(int sockfd, const void *buf, size_t len, int flags);
```

```
ssize_t sendto(int sockfd, const void *buf, size_t len, int flags,  
               const struct sockaddr *dest_addr, socklen_t addrlen);
```

- `send()` is used in connection-oriented sockets (TCP)
- `sendto()` is used in non-connection-oriented sockets (UDP)
- `send(sockfd, buf, len, flags);` is equivalent to `sendto(sockfd, buf, len, flags, NULL, 0);`
- `send(sockfd, buf, len, 0);` is equivalent to `write(sockfd, buf, len);`

Server: step 5 example using send()

```
int fd = socket(AF_INET, SOCK_STREAM, 0);  
...  
int client_fd = accept(fd, (struct sockaddr *)& client_addr,  
                      (socklen_t*)&addrlen);  
...  
  
char msg[] = "hello, world";  
ssize_t r = send(client_fd, msg, strlen(msg), 0);  
if(r < 0) {  
    printf("Error sending message");  
    close(client_fd);  
    exit(0);  
}
```

Server: step 5

- Send and **receive** data

```
ssize_t recv(int sockfd, void *buf, size_t len, int flags);
```

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags,  
                 struct sockaddr *src_addr, socklen_t *addrlen);
```

- *sockfd* is the socket file descriptor (returned by socket())
 - *buf* is a pointer to buffer to be sent
 - *len* is the length of buffer to be sent
 - *flags* is bitwise OR of zero or more options
 - *dest_addr* is a pointer to the structure struct sockaddr which will contain the details of the peer socket
 - *addrlen* is a pointer to the length of what dest_addr points to
-
- If successful, returns number of characters received, otherwise, returns -1
 - If peer socket is shutdown/closed, will return 0

Server: step 5

- Send and **receive** data

```
ssize_t recv(int sockfd, void *buf, size_t len, int flags);
```

```
ssize_t recvfrom(int sockfd, void *buf, size_t len, int flags,  
                 struct sockaddr *src_addr, socklen_t *addrlen);
```

- `recv()` is used in connection-oriented sockets (TCP)
- `recvfrom()` is used in non-connection-oriented sockets (UDP)
- `recv(sockfd, buf, len, flags);` is equivalent to `recvfrom(sockfd, buf, len, flags, NULL, 0);`
- `recv(sockfd, buf, len, 0);` is equivalent to `read(sockfd, buf, len);`

Server: step 5 example using recv()

```
int fd = socket(AF_INET, SOCK_STREAM, 0);
...
int client_fd = accept(fd, (struct sockaddr *)& client_addr,
                      (socklen_t*)&addrlen);
...

char incoming[100];
ssize_t r = recv(client_fd, incoming, 100, 0);
if(r <= 0) {
    printf("Error sending message");
    close(client_fd);
    exit(0);
}
// Do something with receiving message
printf("Received message: %s", incoming);
```

Client: step 1

- Create a socket with the `socket()` system call
- Same as server step 1

Client: step 2

- Connect the socket to the address of the server using the `connect()` system call
 - This step is only required for connection-oriented sockets (TCP)

```
int connect(int sockfd, const struct sockaddr *addr,  
socklen_t addrlen);
```

- *sockfd* is the socket file descriptor (returned by `socket()`)
- *addr* is a pointer to the structure `struct sockaddr` which will contain the details of the server socket
- *addrlen* is a pointer to the length of what *addr* points to
- If successful, returns 0, otherwise, returns -1

Client: step 3

- Send and receive data
- Same as server step 5

Closing a socket

- Socket must be closed after its use

```
int shutdown(int sockfd, int how);
```

```
int close(int sockfd);
```

- *sockfd* is the socket file descriptor (returned by `socket()`)
- *how* can either be `SHUT_RD` (further receptions disallowed), `SHUT_WR` (further transmissions disallowed), or `SHUT_RDWR` (further receptions and transmissions disallowed)
- If successful, returns 0, otherwise, returns -1

Next lecture

- More on socket programming